

1 Overview

- Week 1-3: Classical AI, search algorithms
 1. Uninformed search
 2. Local search: hill climbing
 3. Informaed search: A*
 4. Adversarial search Minimax
- Week 4-7: Classical ML
 1. Decision trees
 2. Linear/Logistic regression
 3. Kernels and support vector machines
 4. "Classical" unsuperivese learning
- Week 10-12: Modern ML
 1. Neural networks
 2. Deep learning
 3. Sequential data
- Week 13: Misc.

2 AI: Computers Trying to Behave Like Humans

- **PEAS Framework:**
 - **Performance measure:** define “goodness” of a solution
 - **Environment:** define what the agent can and cannot do
 - **Actuators:** outputs
 - **Sensors:** inputs
- Agent function is sufficient.
- Common agent structures (to define an AI agent):
 - Reflex
 - Goal-based
 - Utility-based
 - Learning
 - (Others possible; can mix and match!)
- Exploration vs exploitation

3 Problem Statement

fully observable \wedge deterministic \wedge static \wedge discrete \implies only need to observe once

To solve a prob using search:

- A goal or a set of goals
- a model of the environment
- a search algorithm

goal formulation \rightarrow problem formulation \rightarrow search \rightarrow execute

1. goal formulation

2. problem formulation, eg. path finding

- states: nodes representation invariant:: abstract states should correspond to concrete states
- initial state: starting node
- goal states/test: dest node
Goal test: define the goal using a function *is_goal*
- actions: move along an edge :: $|actions(state)| \leq branching_factor$
- transition model: $f(curr_state, action) \implies next_state$
- action cost function: see edges

3. Important facts:

- Representation Invariant: ensure that the abstract states correspond to concrete states
- Goal Test: Goal defined via a function *is_goal*
- Action: a set of $action(state)$, $|actions(state)| \leq branching_factor$
- Transition model: $f(curr_state, action) \implies next_state$

Search

Uninformed search

No information that could guide the search: no clue how good a state is

```
create frontier
// create visited // with vsited memory
insert Node(initial_state) to frontier
while frontier is not empty:
    node = frontier.pop()
    if node.state is goal:
        return solution
// if node.state in visited: // with vsited memory
//     continue
// visited.add(state)
for action in actions(node.state):
    next_state = transition(node.state, action)
    frontier.add(Node(next_state))
return failure
```

Different subvariant of tree search uses different DS for the frontier.

Search Type	Data Structure for Frontier
BFS	Queue
DFS	Stack
UCS (Uniform-cost Search)	Priority Queue

Depth limited search

limit the search to depth l

backtrack when the limit is hit.

time complexity: exponential to search depth

space complexity: size of the frontier

```
create frontier
tier = 0
insert Node(initial_state) to frontier
while (!empty(frontier)) && (tier <= limit):
    node = frontier.pop()
    tier++
    if node.state is goal:
        return solution
    for action in actions(node.state):
        next_state = transition(node.state, action)
        frontier.add(Node(next_state))
return failure
```

Iterative deepening search

search with depth from 0 to ∞

return soln when found. Both complete

```
create frontier
tier = 0
insert Node(initial_state) to frontier
while (!empty(frontier)) && (tier <= limit):
    node = frontier.pop()
    tier++
    if node.state is goal:
        return solution
    for action in actions(node.state):
        next_state = transition(node.state, action)
        frontier.add(Node(next_state))
return failure
```

Summary

Name	Time Complexity*	Space Complexity*	Complete?	Optimal?
Breadth-first Search	Exponential	Exponential	Yes	Yes
Uniform-cost Search	Exponential	Exponential	Yes	Yes
Depth-first Search	Exponential	Polynomial	No#	No
Depth-limited Search	Exponential	Polynomial**	No**	No**
Iterative Deepening Search	Exponential	Exponential**	Yes	Yes

Not complete if not tracking visited nodes, search may stuck in loop before visiting all nodes.

* In terms of some notion of depth/tier

** If used with DFS

4 Local Search

Systematic search: typically complete and optimal under certain constraints. However intractable sometimes
Local search: typically incomplete and suboptimal, but has anytime property, ie. longer time -> better solution. Able to provide good enough solution under reasonable amount of time.

4.1

1. Start at random position in the state space
2. iteratively move from a state to another neighbouring state via perturbation or construction
3. solution is the final state

State space: all possible configuration (1)

Search space: a subset of state space that will be explored (2)

4.1.1 Perturbation search

- Search space: complete candidate solutions
- search step: modification of one or more solution

For example: swap a path with another path

4.1.2 Constructive search

- partial candidate solution
- extension of one or more solution

For example: path finding

4.2

goal formulation -> problem formulation -> search -> execute

1. goal formulation
2. problem formulation, eg. path finding
 - states: nodes representation invariant:: abstract states MAYNOT directly correspond to concrete states
 - initial state: starting node, a candidate solution
 - goal states/test: dest node [optional]
Goal test: define the goal using a function $is_goal\ f(curr_state, action) \implies next_state$
 - Successor function: a function that generates neighbouring states by applying modifications from the current state. This defines the local search space

4.3 Evaluation function

A math function that assess the quality or desirability of the solution.

Some solutions may be unacceptable but there are some less bad than the others.

4.4 Hill Climbing/ Greedy Local Search

```
curr_state = init_state
while 1:
    best_succ = best(successor(curr_state))
    if (eval(best) <= eval(curr_state))
        return curr_state
    curr_state = best_succ
```

4.5 State space landscape

- Global max:
- Local max:
- shoulder:

5 Adversarial search

5.1 Classical adversarial games

- Fully observable
- Deterministic
- discrete
- No infinite run
- 2-player zero-sum
- turn taking

terms

- Player: agent
- Turn:
- Move
- End state
- winning conditon
-

5.2 Problem formulation in adversarial search

- states: nodes representation invariant:: abstract states MAYNOT directly correspond to concrete states
- initial state: starting node, a candidate solution
- Terminal State: state the outcome of the game when it terminates
- Utility function: output the value of a state from the perspetive of our agent

5.3 Minimax

In the view of A, A try to maximize the outcome of the game, B will try to minimize A's outcome, as the gam is zero sum

- $\text{expand}(\text{state}) \Rightarrow [a]$

```
max_value(state):
    if is_terminal(state): return utility(state)
    v = -∞
    for next_state in expand(state):
        v = max(v, min value for player A in next_state)
    return v

min_value(state):
    if is_terminal(state): return utility(state)
    v = ∞
    for next_state in expand(state):
        v = min(v, max value for player A in next_state)
    return v

minimax(state):
    v = max_value(state)
    return action in expand(state) with value v
```

5.4 Alpha-beta prunning

From the viewpoint of the MAX player:

- α : the value of the best choice for the MAX player so far
- β : the value of the best choice for the MIN player so far

An optimized version of the Minimax algorithm using pruning:

```
max_value(state, α, β):
    if is_terminal(state): return utility(state)
    v = -∞
    for next_state in expand(state):
        v = max(v, min(v, α, β))
    return v

min_value(state, α, β):
    if is_terminal(state): return utility(state)
    v = ∞
    for next_state in expand(state):
        v = min(v, max(v, α, β))
    return v

alpha-beta search(state):
    v = max_value(state, -∞, ∞) // initialized α to be -∞, β to ∞
    return action in expand(state) with value v
```

6 Learning agent

For problems that the function are difficult to specify, solutions re intractable to compute in general. Typically episodic,

$$DL \subset ML \subset AL$$

6.1 supervised

Learn the mapping input, feedback given \rightarrow output, given a dataset, it minimizes the difference between the prediction and the provided correct ans using a learnign algorithm e.g. image ientification

1. Train phase: try minimizing the diff between pred and correct ans given using the training set, resulting in a trainign agen function, known as the model/hypothesis
2. Testing/evaluation phase: using a test set to measure the performance of the model. The performance on unseen data measures the generalization of the model

Task

- Classification: to predict discrete labels or catagories on the input features
- Regression: predict continuous numetical value based on input features

Dataset

$$D = \bigcup_{[1,n]} \{(x^{(i)}, y^{(i)})\}$$

True data generation function

$$y = f^*(x) + \epsilon$$

where $f^*(x)$ is true but unknow, which generates the label from the input features; ϵ is some noise or error term, which account for the randomness or imperfection in the date generating process.
the goal is to find a function that best approximately $f^*(x)$

Hypothesis class

THE set of models or functions that maps from inputs to outputs $h : X \implies Y$ that can be learned by a learing algorithm. Each element of the hypo clas $h \in \mathcal{H}$

LEarnign algorithm

$$A(D_{train}, H_{hypo}) = h(x), h \in \mathcal{H} \approx f^*(x)$$

PERformance measure

$$h(x) \approx f^*(x)$$

$$PM(D_{test}, h \in H_{hypo}) \mapsto$$

Try the hypothesis h on a new set of examples (test data)

Regression: error

$$\text{Absolute error} = |\hat{y} - y| \quad (3)$$

$$\text{Squared error} = (\hat{y} - y)^2 \quad (4)$$

Mean squared error

$$MSE = \frac{1}{N} \sum_{i=0}^N (\hat{y}^{(i)} - y^{(i)})^2$$

Mean absolute error

$$MAE = \frac{1}{N} \sum_{i=0}^N |\hat{y}^{(i)} - y^{(i)}|$$

where

$$\hat{y}^{(i)} = h(x^{(i)})$$

Accuracy:

$$A = \frac{1}{N} \sum_{i=1}^n \mathbb{1}_{\hat{y}^{(i)} = y^{(i)}}$$

confusion matrix

True positive, false positive, false negative, true negative (2, 1, 3, 4 quadrant)

$$\text{Accuracy} = \frac{TP + TN}{TP + FN + FP + TN}$$

$$\text{Precision} = \frac{TP}{(TP + FP)}$$

maximize if FP is costly

$$\text{Recall} = \frac{TP}{(TP + FN)}$$

maximize recall if FN is dangerous

$$F1 = \frac{2}{\frac{1}{P} + \frac{1}{R}}$$

6.2 Decision Tree

Greedy, top-down, recursive, use

```
DTL(examples, attributes, default):
    if (examples == ∅): return default
    if (∀e ∈ examples, e has the same classification c): return c
    if (attributes == ∅): return mode(examples)

    best = choose_attribute(attributes, examples)
    tree = new decision tree with root test best
    for vi of best do:
        examples_i = {e | e ∈ examples, e.best = vi}
        subtree = DTL(examples_i, attributes \ best, mode(examples_i))
        tree.add(vi: subtree)
```

$$f : [\text{attribute vector}] \mapsto \text{Boolean}$$

Basically nested if-else

Expressiveness

Decision trees can express any function of the input attributes. Trivially, Consistent training set \rightarrow Consistent decision tree, but unlikely to generalize to new examples

Each row in the truth table is represented as a path in the decision tree

Size of the hypothesis class

For n boolean attributes, there are n boolean functions, n distinct truth tables each with 2^n rows $\rightarrow 2^{2^n}$ decision trees

Informativeness

Ideally we want to select an attribute that splits the examples into all positive or all negative.

Entropy: The higher the more random, thus less informative

$$I(P(v_1) \dots P(v_k)) = - \sum_{i=1}^k P(v_i) \log_2 P(v_i)$$

For data set contains boolean outputs,

$$I(P(+), P(-)) = -\frac{p}{p+n} \log_2 \frac{p}{p+n} - \frac{n}{p+n} \log_2 \frac{n}{p+n}$$

where $0 \leq \mathbb{R}_I \leq 1$. However for non-binary variables the entropy can be greater than 1

Information gain

Information gain = entropy of this node - entropy of children nodes

$$IG(A) = I\left(\frac{p}{p+n}, \frac{n}{p+n}\right) - remainder(A)$$

$$remainder(A) = \sum_{i=1}^v \frac{p_i + n_i}{p+n} I\left(\frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i}\right)$$

Decision tree pruning

- By sample size
- by max search depth

6.3 unsupervised

Find pattern. No feedback given. e.g, group images by cluster

6.4 Reinforcement

Trial and error, reward given based on observation and action. eg. chess